

¹ QuantumUQ: Uncertainty Quantification for Quantum Machine Learning with PennyLane and Qiskit

³ Ferhat Catak¹

⁴ 1 University of Stavanger, Norway

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review ↗](#)
- [Repository ↗](#)
- [Archive ↗](#)

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))⁷.

⁵ Summary

⁶ Quantum machine learning models can produce predictions that look confident even when they
⁷ are unreliable. This can happen because quantum computers estimate outputs by sampling
⁸ measurements (“shots”), because near-term devices are noisy, and because training quantum
⁹ models can be unstable. **QuantumUQ** is an open-source Python library that helps researchers
¹⁰ measure and compare predictive uncertainty in quantum machine learning workflows across
¹¹ two widely used ecosystems: **PennyLane** and **Qiskit**.

¹² QuantumUQ provides tools to estimate shot-based variability, approximate model (epistemic)
¹³ uncertainty using ensembles, profile prediction stability across different shot budgets, and
¹⁴ evaluate reliability using calibration metrics such as reliability diagrams and expected calibration
¹⁵ error (ECE). The library targets researchers who need reproducible, backend-agnostic uncertainty
¹⁶ analysis for variational quantum classifiers and related QML models.

Statement of need

¹⁸ Uncertainty quantification (UQ) is a key element of trustworthy machine learning because it
¹⁹ helps detect when a model’s output should not be trusted. In quantum machine learning (QML),
²⁰ uncertainty has additional sources that are intrinsic to quantum computation: measurement
²¹ outcomes are stochastic and expectation values are estimated with a finite number of shots,
²² leading to sampling variance that depends directly on the measurement budget ([Nielsen & Chuang, 2010](#)). Furthermore, NISQ-era devices introduce hardware noise that perturbs
²³ measurement statistics ([Preskill, 2018](#)). Variational quantum circuits can also be difficult
²⁴ to train due to non-convex optimization and barren plateaus ([McClean et al., 2018](#)), and
²⁵ uncertainty-aware approaches are increasingly important for reliable evaluation ([Kendall & Gal, 2017](#)).
²⁶

²⁸ While QML frameworks provide powerful tools to construct and train circuits, they do not
²⁹ provide a unified, research-focused interface for uncertainty estimation, shot-budget profiling,
³⁰ and calibration analysis across backends. As a result, QML studies often report accuracy
³¹ without systematically characterizing reliability. QuantumUQ fills this gap by providing a
³² backend-agnostic UQ layer aimed at QML researchers who need consistent uncertainty metrics,
³³ reproducible experiments, and explicit resource-reliability trade-off analysis.

³⁴ State of the field

³⁵ PennyLane and Qiskit support circuit construction, simulators/hardware execution, and
³⁶ variational training. However, uncertainty-related evaluations are typically implemented ad
³⁷ hoc and are not standardized across frameworks. Classical uncertainty libraries and calibration
³⁸ toolkits are widely used in conventional deep learning, but they do not model shot-driven

39 uncertainty or measurement-budget trade-offs in QML. Calibration for probabilistic classifiers
40 is well studied in classical ML ([Guo et al., 2017](#)), yet QML papers rarely include systematic
41 calibration diagnostics.

42 QuantumUQ contributes a QML-specific layer that standardizes shot-based uncertainty
43 estimation, ensemble-based epistemic uncertainty ([Lakshminarayanan et al., 2017](#)), stability
44 profiling across shot budgets, and calibration metrics. It also provides adapters that make
45 these evaluations consistent across PennyLane and Qiskit. This design enables reproducible
46 research and fair comparisons across backends.

47 Software design

48 QuantumUQ is organized around (i) backend adapters, (ii) uncertainty methods, (iii) a
49 predictive distribution representation, and (iv) evaluation metrics.

50 Backend adapters normalize execution across PennyLane and Qiskit, enabling repeated circuit
51 evaluations at specified shot budgets. A `PredictiveDistribution` abstraction represents
52 repeated stochastic predictions and exposes common uncertainty summaries (mean, variance)
53 and reliability metrics (entropy and calibration).

54 QuantumUQ implements two core uncertainty estimators: shot-based estimators that quantify
55 measurement-induced variability, and ensemble-based estimators that approximate epistemic
56 uncertainty by training multiple independent models ([Lakshminarayanan et al., 2017](#)). The
57 library also includes stability profiling utilities that characterize how predictions change as shot
58 count increases, supporting principled selection of measurement budgets.

59 Design trade-offs are explicit: higher shot counts reduce sampling variance but increase runtime,
60 while ensemble methods increase cost linearly in the number of trained models. QuantumUQ
61 makes these trade-offs measurable and reproducible for QML workflows.

62 Example Outputs

63 QuantumUQ includes utilities for calibration analysis and measurement-budget profiling.
64 [Figure 1](#) shows a reliability diagram generated using the library, illustrating mild overconfidence
65 in high-confidence regions. [Figure 2](#) shows predictive entropy as a function of the number
66 of measurement shots for both PennyLane and Qiskit backends, demonstrating the expected
67 decrease in uncertainty as the shot budget increases.

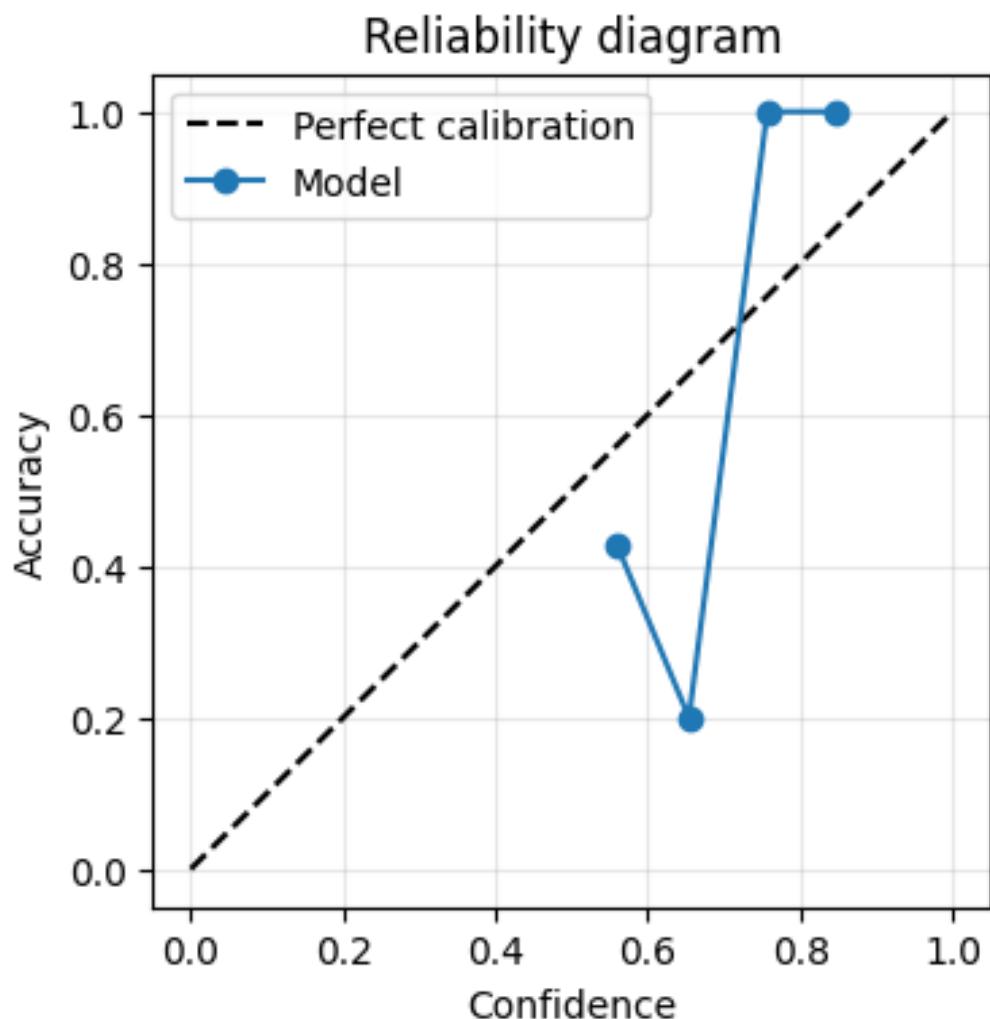


Figure 1: Reliability diagram illustrating calibration behavior of a variational quantum classifier.

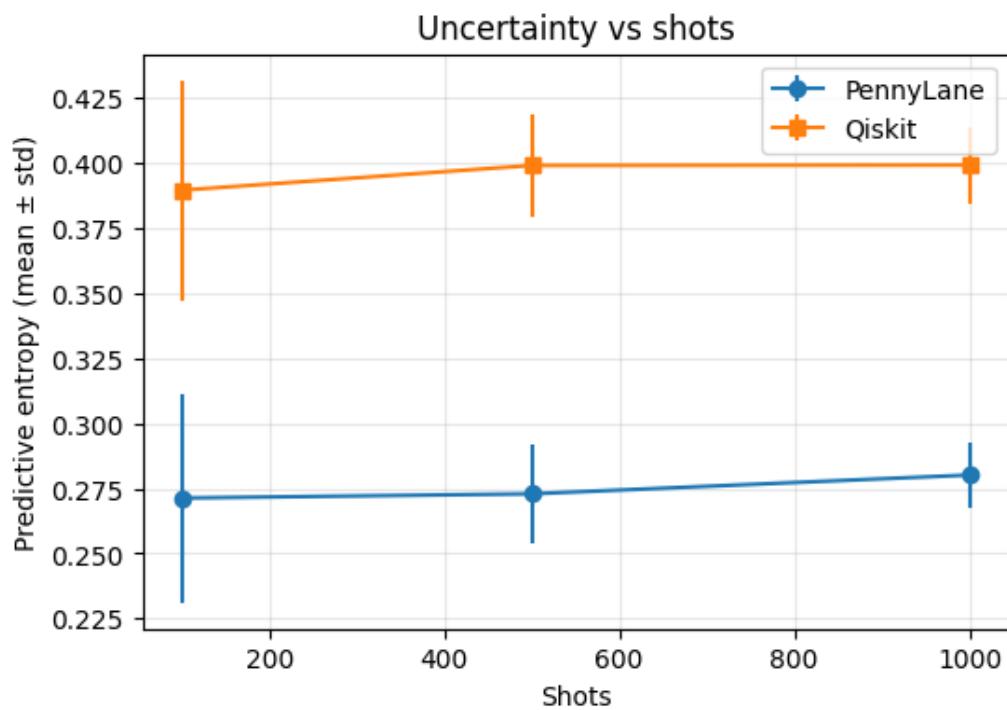


Figure 2: Predictive entropy as a function of measurement shots for PennyLane and Qiskit backends.

68 Research impact statement

69 QuantumUQ is distributed as an installable Python package and includes documentation and
70 executable notebooks demonstrating usage with both PennyLane and Qiskit. The repository
71 provides reproducible materials for generating calibration plots and uncertainty-versus-shots
72 curves for variational quantum classifiers (Schuld et al., 2020). The package supports research
73 questions central to reliability in QML: how many shots are required for stable predictions,
74 where models are overconfident, and how uncertainty behavior differs across backends.
75 The software is designed to support near-term research impact by enabling consistent evaluation
76 across QML frameworks and by providing reusable, reproducible building blocks for trustworthy
77 quantum machine learning experiments.

78 AI usage disclosure

79 (Choose one.) - No generative AI tools were used in the software development, documentation,
80 or paper writing. - Generative AI tools were used to assist with drafting and editing parts of
81 the manuscript and/or documentation. All AI-assisted content was reviewed and edited by
82 the author, and correctness was verified via unit tests and rerunning the provided notebooks
83 end-to-end.

84 Acknowledgements

85 (Add funding, institutional support, compute resources, or project acknowledgements here.)

86 References

- 87 Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural
88 networks. *Proceedings of the 34th International Conference on Machine Learning - Volume*
89 *70*, 1321–1330.
- 90 Kendall, A., & Gal, Y. (2017). What Uncertainties Do We Need in Bayesian Deep Learning
91 for Computer Vision? *arXiv e-Prints*, arXiv:1703.04977. <https://doi.org/10.48550/arXiv.1703.04977>
- 93 Lakshminarayanan, B., Pritzel, A., & Blundell, C. (2017). Simple and scalable predictive
94 uncertainty estimation using deep ensembles. *Advances in Neural Information Processing*
95 *Systems, 30*.
- 96 McClean, J. R., Boixo, S., Smelyanskiy, V. N., Babbush, R., & Neven, H. (2018). Barren
97 plateaus in quantum neural network training landscapes. *Nature Communications*, *9*(1),
98 4812.
- 99 Nielsen, M. A., & Chuang, I. L. (2010). *Quantum computation and quantum information*.
100 Cambridge university press.
- 101 Preskill, J. (2018). Quantum computing in the NISQ era and beyond. *Quantum*, *2*, 79.
- 102 Schuld, M., Bocharov, A., Svore, K. M., & Wiebe, N. (2020). Circuit-centric quantum
103 classifiers. *Phys. Rev. A*, *101*, 032308. <https://doi.org/10.1103/PhysRevA.101.032308>

DRAFT